



**FACULTY
OF MATHEMATICS
AND PHYSICS**
Charles University

BACHELOR THESIS

Devyanshu Koirala

Artificial Intelligence for Quoridor game

Department of Software and Computer Science Education (KSVI)

Supervisor of the bachelor thesis: Mgr. Klára Pešková, Ph.D.

Study programme: Computer Science

Study branch: Artificial Intelligence (AI)

Prague 2023

I declare that I carried out this bachelor thesis independently, and only with the cited sources, literature and other professional sources. It has not been used to obtain another or the same degree.

I understand that my work relates to the rights and obligations under the Act No. 121/2000 Sb., the Copyright Act, as amended, in particular the fact that the Charles University has the right to conclude a license agreement on the use of this work as a school work pursuant to Section 60 subsection 1 of the Copyright Act.

In date
Author's signature

Dedication.

Title: Artificial Intelligence for Quoridor game

Author: Devyanshu Koirala

Department: Department of Software and Computer Science Education (KSVI)

Supervisor: Mgr. Klára Pešková, Ph.D., KSVI

Abstract: Quoridor presents a challenging terrain for strategic decision-making, making it a suitable testing ground for various AI algorithms. This thesis explores the implementation and evaluation of three distinct AI algorithms, namely A*, Minimax with Alpha-Beta Pruning, and Monte Carlo Tree Search (MCTS), within the realm of Quoridor gameplay.

The research begins with a comprehensive overview of the Quoridor game, its rules and strategies. Subsequently, we delve into the theoretical foundations and practical implementation details of the aforementioned AI algorithms and conduct a thorough evaluation in an effort to determine the best one in this context.

Keywords: Quoridor Artificial Intelligence AI Board Game

Contents

Introduction	2
0.1 Acknowledged Works	3
1 Background	4
1.1 Zero-sum game	4
1.2 Minimax algorithm	4
2 Description	6
2.1 Notation	6
2.2 Rules	8
2.2.1 Wall placement rules	8
2.2.2 Player movement rules	9
3 Game Analysis	10
3.1 Classification of Quoridor	10
3.1.1 Discrete	10
3.1.2 Deterministic	10
3.1.3 Zero-sum	10
3.1.4 Perfect Information	10
3.2 Game-Tree	11
3.2.1 Branching Factor	11
3.3 State space complexity	14
3.4 Game tree complexity	14
3.5 Comparison with other games	15
Conclusion	16
Bibliography	17
List of Figures	18
List of Tables	19
List of Abbreviations	20
A Attachments	21
A.1 First Attachment	21

Introduction

Artificial Intelligence (AI), over the past decade, is becoming an integral part of many elements of modern world. In fact, in the recent years, it has become an the main point of evolution and revolution in many of the technologies and industries. From assisted or autonomous driving [Ma et al., 2020], chat bots Wu et al. [2023] to gaming Skinner and Walmsley [2019], AI has been a major revelation in evolving the existing technologies to generating new industry space with the plethora of new use cases.

AI has brought a revolutionary transformation to the gaming world, pushing the boundaries of what’s achievable in both single and multiplayer gaming experiences. In recent years, AI has taken center stage with its remarkable accomplishments in mastering age-old games such as Chess, Go, and many others. AI-driven games now offer users the opportunity to hone and enhance their skills, providing varying difficulty levels and offering optimal moves to guide players through each step if desired.

Strategy games are a genre of gaming that require planning, often involving various tactics, decision making and execution while under various constraints (e.g., resources). Some of the examples of popular strategy games are Chess, Go, Shogi, Starcraft, etc. Strategy games are unique compared to many other genres of games in the sense that it requires a selection of an optimal move among a subset of sub-optimal moves based on a plan. In many scenarios, the size of the subset of sub-optimal moves depends on exploration of the game space, simulation and prediction of sequence of the player’s and the opponent’s moves all while managing certain resources efficiently.

AI, due to its suitability towards solving problems involving complex decision making, taking into account sequences of the player’s and the opponents possible moves (e.g., game space), outcome prediction or position evaluation all while considering the available resources, have been deemed particularly effective in playing the strategy games. The history of the use of AI in strategy games goes back to the last few decades. One of the initial marked impact of AI in the strategy gaming came in 1997 when IBM’s Deep Blue Campbell et al. [2002] defeated World Chess Champion Garry Kasparov. The influence of AI in strategy games was more prominent with the rise and involvement of AI into the real-time strategy (RTS) games such as Warcraft and StarCraft Robertson and Watson [2014] and implementation of Monte-Carlo tree search (MCTS) in the strategy games such as Go Huang et al. [2011]. Recently, DeepMind’s Alpha Go for Go, Alpha Zero for Chess Silver et al. [2017] and AlphaStar for StarCraft Team [2019] have widened the gap between the AI and human intelligence even further.

Designed by Mirko Marchesi, Quoridor stands out as an engaging strategic board game that is played between two or four players. The game is played on a square grid board where the objective of this game is for each player to move their pawn to the opposite side of the board. This game introduces a fascinating twist where a player, in addition to trying to move their pawn through the square grid, additionally have an option to place walls on the grid locations strategically to obstruct the opponent’s path. This strategy compels the player to think of their traversal strategy while predicting the opponents strategy as well.

Despite its seemingly simple rules, Quoridor demands a unique blend of strategic foresight and the ability to anticipate the moves of opponents and outmaneuver the opponent.

Compared to some other strategy games such as Chess and Go, Quoridor has not been extensively studied in the literature. In Lisa Glendenning [2002], the author developed an agent for playing Quoridor using genetic algorithm to optimize the weights. The authors in Mertens [2006] study the complexity of the algorithm also develop a Quoridor playing agent based on Minimax algorithm. The authors conclude that Quoridor has a similar state-space and game tree complexity as that of the games such as Chess. Likewise, the authors in Brenner [2015] developed an MCTS approach for Quoridor. Recently, in Iwanaga and Sakamoto [2022], the authors performed an analysis of the game for a miniature 5 by 5 board.

The primary objective of this thesis is to construct a well-structured framework and user-friendly interfaces that seamlessly integrate AI algorithms into the Quoridor game. The development of AI algorithms customized to Quoridor’s unique rule set will not only enhance our understanding of the game’s intricate nuances but also facilitate the creation of an intuitive interface for simulating these AI agents. Furthermore, a comprehensive evaluation will be conducted to identify the top-performing AI agent.

In addition to this, the project will encompass the creation of a user-friendly interface that empowers players to engage with an AI opponent of their choice, thereby bolstering the game’s accessibility and inclusivity.

0.1 Acknowledged Works

Quoridor, being a widely popular game, has attracted a fair number of attention from the research community, resulting in successful AI agent developments. Some of the notable works include:

- **Mastering Quoridor [Lisa Glendenning, 2002]** The writer implements and assesses various algorithms like Negamax, Alpha-beta negamax among others. Additionally, they utilized a genetic algorithm to refine the weights within a linear weighted evaluation function, employing 10 distinct features suggested by the author, some of which include player’s position towards their goal side, the opponent’s position towards their respective goal, the remaining count of walls available to the player, etc.

In sharp contrast to the aforementioned paper, our thesis takes a distinctly different path by delving deeply into the architectural aspects of AI. Our approach emphasizes abstraction to the greatest extent possible, with an eye on facilitating seamless integration into a broad spectrum of games. We prioritize creating an interface that is adaptable to diverse game environments, setting our research apart from the game-specific focus of the prior works.

1. Background

In this chapter, we give a brief overview of different AI techniques that have been considered for implementing the agent for Quoridor including the minimax algorithm.

1.1 Zero-sum game

A zero-sum game is defined as a game where an advantage to one side means an equal loss to the opponent, resulting in the sum of zero. In zero sum games such as chess, tic-tac-toe, an advantage or a win for a team would mean a disadvantage or a loss for another team.

An example of a zero sum game is poker where, for example, 5 people buy in with a total of 50 euros each, making the table total of 250 euros. In the end, despite some winners or losers in the game, the total will be distributed among the players. In this game, there will be a zero net gain resulting in a zero sum game.

Similarly, in chess, the evaluation of a given position on the board may be slightly advantageous to white. This will always imply that the position is slightly disadvantageous for the player playing the black pieces. A loss of a rook for black would mean, in some cases, an equivalent loss worth 5 pawns. This would imply an equivalent 5 pawn gain for the opponent playing with the white pieces.

1.2 Minimax algorithm

Minimax algorithm, first proven by John von Neumann in 1928 in his paper *Zur Theorie Der Gesellschaftsspiele*, is a very popular algorithm employed in many decision-making scenarios for e.g., in decision theory, game theory and even philosophy. As suggested by the name minimax, the idea of the algorithm is to minimize the player's loss when the opponent makes a decision that gives the player the maximum loss. This algorithm has been implemented in many two-player zero-sum strategy games such as chess and tic-tac-toe

Mathematically, a minimax algorithm can be defined as the following:

$$\bar{v}_i = \min_{a_{-i}} \max_{a_i} v_i(a_i, a_{-i}) \quad (1.1)$$

where,

$$i = \text{index of the player of interest} \quad (1.2)$$

$$-i = \text{index of the opponent(s)} \quad (1.3)$$

$$a_i = \text{action of the player of interest} \quad (1.4)$$

$$a_{-i} = \text{action of the opponents} \quad (1.5)$$

$$v_i = \text{value function of player } i \quad (1.6)$$

$$\bar{v}_i = \text{minimax value of the player of interest} \quad (1.7)$$

The value represented by $v_i(a_i, a_{-i})$ defines the initial set of values or outcomes of the game depending on various actions of the player i and player $-i$. We

first marginalize away the set of actions of the the player (e.g., a_i) from the set of possible outcomes by maximizing the value v_i over the set of every possible action of the player. The implication here is that we are evaluating all the possible actions of the player and selecting the one that maximizes the value function. We then choose an action a_{-i} from the set of all possible actions of the opponent that minimizes the maximum value function.

In a two-player zero sum game, this concept translates to the following interpretation. Given a two person zero sum game with finite actions of the players, there exists a value V such that given an opponent's strategy, the value to the player of interest is V . This implies that given the player of interest's chooses the strategy, the opponent's value is $-V$ making the sum zero.

Below in an example of a minimax algorithm in zero sum game:

- Consider two players A and B.
- The players simultaneously choose a number (an integer) between 1 and 5.
- The payoff is the difference between the two chosen numbers. For example, if Player A chooses 5 and player B chooses 1, Player B has to give $5 - 1 = 4$ to Player A.
- This can be exemplified with the payoff matrix below:

	B chooses 1	B chooses 2	B chooses 3	B chooses 4	B chooses 5
A chooses 1	0	-1	-2	-3	-4
A chooses 2	1	0	-1	-2	-3
A chooses 3	2	1	0	-1	-2
A chooses 4	3	2	1	0	-1
A chooses 5	4	3	2	1	0

Table 1.1: Payoff matrix for the described game from the perspective of player A

For the $\max_{a_i} v_i(a_i, a_{-i})$ part, the player A first marginalizes away the actions of the opponent a_{-i} . This results in the following table:

	B chooses 1	B chooses 2	B chooses 3	B chooses 4	B chooses 5
$\max_{a_i} v_i(a_i, a_{-i})$	4	3	2	1	0

Table 1.2: Marginalied value of the payoff matrix after maximizing over the player's actions

We first marginalize away the set of actions of the player (e.g., a_i) from the set of possible outcomes by maximizing the value v_i over the set of every possible action of the player.

We then determine the $\min_{a_{-i}}$ over the set of chosen values and determine that 0 is the minimum when B chooses 5.

We determine that for player A, no matter what B chooses, the optimal strategy to maximize the value to the player i is choosing 5. this is the dominant strategy. This is also known as the dominant strategy when the value of the minimax can be determined only based on player A's actions without depending on the actions of player B.

2. Description

The game begins with an $N \times N$ chess-like board where each square represents a potential position for the game pieces. It contains grooves that runs between the squares where players can place their walls. Each player is represented by a pawn represented by a character label that start on opposite sides, with their pawns located at the center of their respective edge. The primary objective is to be the first player to move their pawn to the row of squares on the opposite side of the game board avoiding any walls deterring its path to the goal while strategically placing walls to deter opponents from reaching their goal squares.

Walls are a fundamental element of the game, allowing players to strategically block their opponent's path and influence the course of the game. Each wall spans across two cells and occupies exactly four cells either horizontally or vertically, effectively creating a barrier between them. At the beginning, each player starts with a set number of walls that they can use during their turn.




Figure 2.1: A 5x5 game board

As seen in figure 2.1, In the beginning, player A starts at cell (2, 0) and player B starts at cell (2, 4). A's goal is to reach the cells highlighted in pink. The gray areas between the cells are grooves for wall placement.

2.1 Notation

There are no official notations for this game. However, some popular ones recognized by the Quoridor community include **Glendenning's Notation** ([Lisa Glendenning, 2002]) and the **Quoridor Strats Notation** ([Quoridor Strats, 2014]).

Let $R = \{a, \dots, i\}$ and $C = \{1, \dots, 9\}$


Both notations follow the same principle of labelling each cell by C_{ij} where $i \in R$ and $j \in C$.

Move M and Wall W are denoted algebraically, where


$$M(C_{ij}) = ij \text{ where } i \in R \text{ and } j \in C,$$

$$W(C_{ij}) = ijD \text{ where } i \in R, j \in C \text{ and } D \in \{h, v\}$$

The difference between the two notations is the way the walls are represented.



(a) Glendenning's Notation: **c3h**



(b) Quoridor Strats Notation: **c4h**

Figure 2.2: Notation differences

As seen in Figure 2.2, the difference lies in the point of reference of the wall. In **Quoridor Strats Notation**, each wall coordinate is represented by the lower-left cell the wall follows along, whereas in **Glendenning's Notation**, each wall is represented by the upper-left cell the wall follows along.

Even though they are widely used, they are very easy to get confused with since they have the same wall representations, and unless specified explicitly, it is difficult to tell which representation is being used.

This is why I have decided to use a different representation for walls. Instead of vertical and horizontal wall with respect to a cell, we define a direction explicitly.


$$W(C_{ij}) = ijD \text{ where } i \in R, j \in C \text{ and } D \in \{N, S, E, W\}$$

Looking back at Figure 2.2a, the walls can now be represented by **c4N**, i.e. a Northern wall from the cell C_{c4} .


2.2 Rules

2.2.1 Wall placement rules

- Walls cannot be placed diagonally.
- A placed wall must not completely block any player's path to victory. Each player must have at least one path to victory (*See figure 2.3*)
- A placed wall cannot intersect any of the previously placed walls.
- Walls cannot be placed along the edges of the board. Walls must be placed to create a barrier for exactly 4 cells
- Every player possesses a limited supply of walls, and once they exhaust these walls, they are unable to place any additional ones. Consequently, the player is only allowed to maneuver their pawn on the board.



(a) Valid game state



(b) Invalid game state

Figure 2.3: Example of an invalid wall placement

The game state represented by Figure 2.3a shows the situation after 5 turns, with it currently being player B's turn to move. Since both **A** and **B** have viable paths to their respective goal rows and all walls have been placed according to the rules (see *Section 2.2.1*), the game state shown in Figure 2.3a is considered valid.

However, player B disrupts the rules by placing the red wall, violating the specified wall-placement rules (see *Section 2.2.1*), consequently rendering the game state represented by Figure 2.3b invalid.

2.2.2 Player movement rules

- Players are allowed to move their pawn one unit at a time in the North, South, East, or West directions during their turn. Diagonal movements are not allowed.
- **Jump**
 - If an opponent is at to the cell a player intends to move to, the player can jump over the opponent provided there is no wall between them or behind the opponent they intend to jump over. In the latter case, the player can jump to a cell on either side of the opponent's cell, given the cell is accessible from the opponent's cell.
 - Players cannot jump over walls
 - A Player is allowed to jump over multiple opponents as long as they adhere to the aforementioned conditions.




Figure 2.4: Examples of possible moves for player B in the game state

3. Game Analysis

In Chapter 2, we briefly explored the rules and gameplay mechanics of Quoridor. As we progress, this chapter aims to deepen our understanding by analyzing Quoridor from a theoretical and computational perspective.

In this chapter will classify Quoridor within the realm of strategic games, examine its game tree, state-space and tree complexity, and explore the implications of these factors on gameplay and artificial intelligence application.

3.1 Classification of Quoridor

Quoridor can be characterized as a discrete, deterministic, zero-sum, sequential, game with perfect information (Lisa Glendenning [2002]) , and therefore, a combinatorial game (Thomas S Ferguson [2020])

3.1.1 Discrete

In every turn, each player has a finite number of moves and wall placements. These are limited by the game state (already placed walls and moved pawns) and the rules of the game. The game-tree of Quoridor has finite number of nodes (e.g 3.1)

3.1.2 Deterministic

Quoridor has no random elements or chance involved in the gameplay. Every outcome and situation is a direct result of the players' decisions and strategies. There's no dice rolling, card drawing, or any other mechanism that introduces randomness.

3.1.3 Zero-sum

In Quoridor, when a player makes a move that brings them closer to winning (like advancing their pawn or placing a wall effectively), it inherently puts the opponent at a disadvantage. Therefore, any positive progress for a player translates into a negative impact for their opponent. This reciprocal relationship of gain and loss between the players is what characterizes Quoridor as a **zero-sum** game.

3.1.4 Perfect Information

Every aspect of its gameplay are completely visible and known to all players at all times. This means that the positions of the pawns and the placements of the walls on the board are always in full view, allowing players to make strategic decisions based on the entire state of the game.

3.2 Game-Tree

A game tree for an **abstract strategy game** is a comprehensive graph representing every possible game states and sequence of moves. The nodes of a game tree represent game states, and the edges represent action/move.

Game trees are integral to the framework of **adversarial search problems**, where they are employed to systematically explore and evaluate the possible outcomes of different strategies, and forecast future states of the game based on current and potential moves.

The following depicts a (partial) game tree for the Quoridor game with a board of size 3x3:




Figure 3.1: A partial game tree for a 3x3 game board.

3.2.1 Branching Factor

The branching factor of a **Game-tree** is the number of child nodes of each node, or in other words, the number of possible moves a player at their turn can make, given the game state.

In figure 3.1, player A makes the first move. A has **3** places to move their pawn to and **8** places to put one of their walls at. So, the root has a **branching factor** of **11**.

The branching factor is greatly influenced by the state of the board in Quoridor, i.e the player locations and walls placed.

As an example, the figure below represents a game states with the maximum and minimum branching factors:

As depicted by Figure 3.2a, player A has 5 possible places to move their pawn to. No walls have been placed so far, so player A can also place one of their walls in any place.




Figure 3.2: Branching factor differences

For a board of size $N \times N$ with no wall(s) placed, $N - 1$ walls can be placed in each row (since each wall occupies 2 cell lengths), and there are $N - 1$ rows for correct horizontal wall placements. Hence, there are $(N - 1)^2$ slots for horizontal wall placements, and since the board is $N \times N$, the total slots for both horizontal and vertical wall placements is given by the equation:

$$2(N - 1)^2 \quad (3.1)$$

Coming back to figure 3.2a, since the board has no walls placed, we now see that A has $5 + 2 * (5 - 1)^2 = 37$ possible moves they can perform, ergo, the branching factor of the game state represented by Figure 3.2a is **37**, which is also the maximum branching factor for board sized 5x5.

However, in Figure 3.2, A has no available slot for wall-placement, and the already present walls block A from moving anywhere except for cell **c3**. Hence, the branching factor for the game state represented by Figure 3.2 is 1.

Average Branching Factor

In sub-section 3.2.1. we saw that the branching factor is not uniform due factors like wall-placements and positioning of the player greatly influencing it.

We, therefore, would like to estimate an average branching factor for boards of different lengths to see if varying board dimension has any affect in the average branching factor.

We already know from equation 3.1 that the maximum branching factor of the game tree is exponential in order of N and from Figure 3.2, we can deduce that the minimum branching factor is 1 (since we can replicate a similar game state for any game state).

To find an estimate of the average branching factor, we use the Algorithm 1

Algorithm 1: Average branching factor

Function AvgBranchingFactor(*agent1, agent2, simulations*):**input** : Two agents and number of simulations**output**: Average of averages branching factorSumOfAverages \leftarrow 0**for** $i \leftarrow 1$ **to** *simulations* **do** State \leftarrow Initialize() GamePossibleMoves \leftarrow 0 GameMovesMade \leftarrow 0 Agents \leftarrow [agent1, agent2] AgentIndex \leftarrow 0 **while** *State is not Terminal* **do** Agent \leftarrow Agents[AgentIndex] AgentIndex \leftarrow (AgentIndex + 1) % 2 GamePossibleMoves \leftarrow GamePossibleMoves +
 Length(State.PossibleMoves()) Move \leftarrow Agent.GetMove(State) State \leftarrow State.Apply(Move) GameMovesMade \leftarrow GameMovesMade + 1 **end** GameAverage \leftarrow GamePossibleMoves / GameMovesMade SumOfAverages \leftarrow SumOfAverages + GameAverage**end****return** SumOfAverages / simulations

We then simulate **1000** games between **Minimax AB** and **Semi-random** agents, each for boards of dimensions 3, 5, 7 and 9, and the results of the average branching factor can be seen in figure 3.3




Figure 3.3: Branching factor for boards of different sizes

From figure 3.3, like with the maximum branching factor, we can see a similar exponential trend of the average branching factors in order of N . Furthermore, based on only the four dimensions and their averages and maximum branch-

ing factors, the average branching factor seems to be almost half the maximum branching factor

The average branching factor for board of dimension 9x9 was about **62**, which is very close to the value proposed by Lisa Glendenning [2002]

3.3 State space complexity

The state space complexity in Quoridor refers to the total number of possible game states that can be reached from the initial game state. This includes all the possible positions of both players' pawns on the board and all the possible configurations of walls. It is extremely difficult to find an exact value because we also need to account for wall placement rules 2.2.1 and pawn movement rules 2.2.2. Because of this, we loosen the regulations to come up with an upper-bound.

For a board of Dimension $N \times N$, the first pawn can be placed in N^2 possible locations. After the first pawn is placed in the board, the second pawn now has $N^2 - 1$ possible locations to be placed into. So, the total number of ways for pawn placement, S_p is given by: Mertens [2006]

$$S_p = N^2(N^2 - 1) \quad (3.2)$$

As for walls, from equation 3.1, we know that there are $2(N - 1)^2$ ways of placing walls on the board, both horizontally and vertically. Since we know that each wall occupies 2 cells, we will assume that placing the wall anywhere in the board takes away the possibility of placing walls at 4 different places, even though this is not particularly true at the edges. So, the total number of wall placement, S_w is given by: Mertens [2006]

$$S_w = \sum_{i=0}^{20} \prod_{j=0}^i (N^2 - 4j) \quad (3.3)$$

And thus, the total state space complexity is given by $S = S_w * S_p$, which, for the standard quoridor board size 9x9 is $3.9905 * 10^{42}$ Mertens [2006]

3.4 Game tree complexity

The game tree complexity refers to the number of possible games that can be played, i.e the total number of states in the game tree. For example, the root node for game tree for the tic-tac-toe game has a branching factor of 9, or, in other words, the first player has 9 possible places to put the 'X' into. The branching factor decreases by 1 as the game progresses, and hence, the game complexity of tic-tac-toe is $9 * 8 * \dots * 1 = 9!$.

In complex games like Quoridor, the game-tree complexity can be estimated by raising the branching factor to a power of the total number of moves by the players, which is given by Mertens [2006]

$$G = 1.7884 * 10^{162} \quad (3.4)$$

3.5 Comparison with other games

	log(state space complexity)	log(game tree complexity)	Branching factor
Tic-tac-toe	3	5	4
Connect-four	13	21	4
Chess	44	123	35
Quoridor	42	162	60
Go	170	360	250

Table 3.1: State space, game tree and branching factor comparison between well-known games

Conclusion

Bibliography

- Matyáš Brenner. Artificial intelligence for quoridor board game. 2015.
- Murray Campbell, A Joseph Hoane Jr, and Feng-hsiung Hsu. Deep blue. *Artificial intelligence*, 134(1-2):57–83, 2002.
- Shih-Chieh Huang, Rémi Coulom, and Shun-Shii Lin. Monte-carlo simulation balancing in practice. In *Computers and Games: 7th International Conference, CG 2010, Kanazawa, Japan, September 24-26, 2010, Revised Selected Papers 7*, pages 81–92. Springer, 2011.
- Takuro Iwanaga and Makoto Sakamoto. Analysis of 5x5 board quoridor. 2022.
- Lisa Glendenning. Mastering Quoridor. 2002. URL <https://api.semanticscholar.org/CorpusID:18564930>.
- Yifang Ma, Zhenyu Wang, Hong Yang, and Lin Yang. Artificial intelligence applications in the development of autonomous vehicles: a survey. *IEEE/CAA Journal of Automatica Sinica*, 7(2):315–329, 2020. doi: 10.1109/JAS.2020.1003021.
- Peter JC Mertens. A quoridor-playing agent. *Bachelor Thesis, Department of Knowledge Engineering, Maastricht University*, 2006.
- Quoridor Strats. The Official Quoridor Notation, Sep 2014. URL <https://quoridorstrats.wordpress.com/notation/>.
- Glen Robertson and Ian Watson. A review of real-time strategy game ai. *Ai Magazine*, 35(4):75–104, 2014.
- David Silver, Thomas Hubert, Julian Schrittwieser, Ioannis Antonoglou, Matthew Lai, Arthur Guez, Marc Lanctot, Laurent Sifre, Dhharshan Kumaran, Thore Graepel, et al. Mastering chess and shogi by self-play with a general reinforcement learning algorithm. *arXiv preprint arXiv:1712.01815*, 2017.
- Geoff Skinner and Toby Walmsley. Artificial intelligence and deep learning in video games a brief review. In *2019 IEEE 4th International Conference on Computer and Communication Systems (ICCCS)*, pages 404–408, 2019. doi: 10.1109/CCOMS.2019.8821783.
- AlphaStar Team. Alphastar: Mastering the real-time strategy game starcraft ii. *DeepMind blog*, 24, 2019.
- Thomas S Ferguson. *A Course In Game Theory*. World Scientific, 2020.
- Tianyu Wu, Shizhu He, Jingping Liu, Siqi Sun, Kang Liu, Qing-Long Han, and Yang Tang. A brief overview of chatgpt: The history, status quo and potential future development. *IEEE/CAA Journal of Automatica Sinica*, 10(5):1122–1136, 2023. doi: 10.1109/JAS.2023.123618.

List of Figures

2.1	A 5x5 game board	6
2.2	Notation differences	7
2.3	Example of an invalid wall placement	8
2.4	Examples of possible moves for player B in the game state	9
3.1	A partial game tree for a 3x3 game board.	11
3.2	Branching factor differences	12
3.3	Branching factor for boards of different sizes	13

List of Tables

1.1	Payoff matrix for the described game from the perspective of player A	5
1.2	Marginalied value of the payoff matrix after maximizing over the player's actions	5
3.1	State space, game tree and branching factor comparison between well-known games	15

List of Abbreviations

A. Attachments

A.1 First Attachment